

The iterated filtering convergence diagnostics plot in figure 1 comes from a 2021 student project investigating COVID-19. The calculation used 10^3 particles. What is the best interpretation?

A: Everything seems to be working fine. There is a clear consensus from the different searches concerning the highest likelihood that can be found. Therefore, the search is doing a good job of maximization. Occasional searches get lost, such as the purple line with a low likelihood, but that is not a problem.

B: The seaches obtain likelihood values spread over thousands of log units. We would like to see consistent convergence within a few log units. We should use more particles and/or more iterations to achieve this.

C: The seaches obtain likelihood values spread over thousands of log units. We would like to see consistent convergence within a few log units. We should compare the best likelihoods obtained with simple statistical models, such as an autoregressive moving average model, to look for evidence of model misspecification.

D: The seaches obtain likelihood values spread over thousands of log units. We would like to see consistent convergence within a few log units. We should look at the effective sample size plot for the best fit we have found yet, to see whether there are problems with the particle filtering.

E: All of B, C, and D.

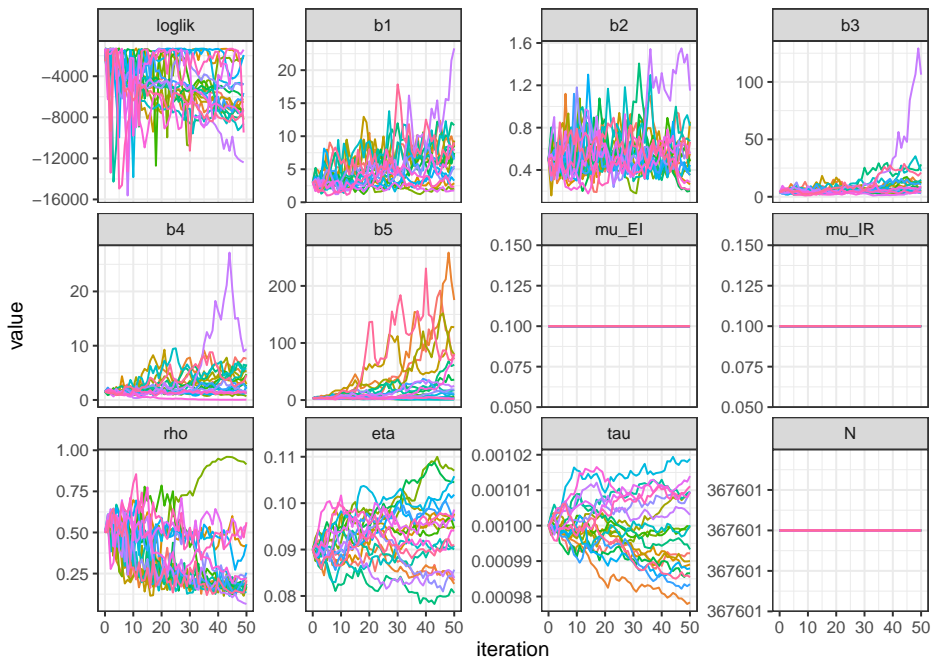


Figure 1: Diagnostic plot for a COVID-19 model